



Heterogeneous Computing

David Luebke





The “New” Moore’s Law

- Computers no longer get faster, just wider
- You **must** re-think your algorithms to be parallel !
- Data-parallel computing is most scalable solution

Enter the GPU



- Massive economies of scale
- Massively parallel



Enter CUDA



- Scalable parallel programming model
- Minimal extensions to familiar C/C++ environment
- Heterogeneous serial-parallel computing

Sound Bite



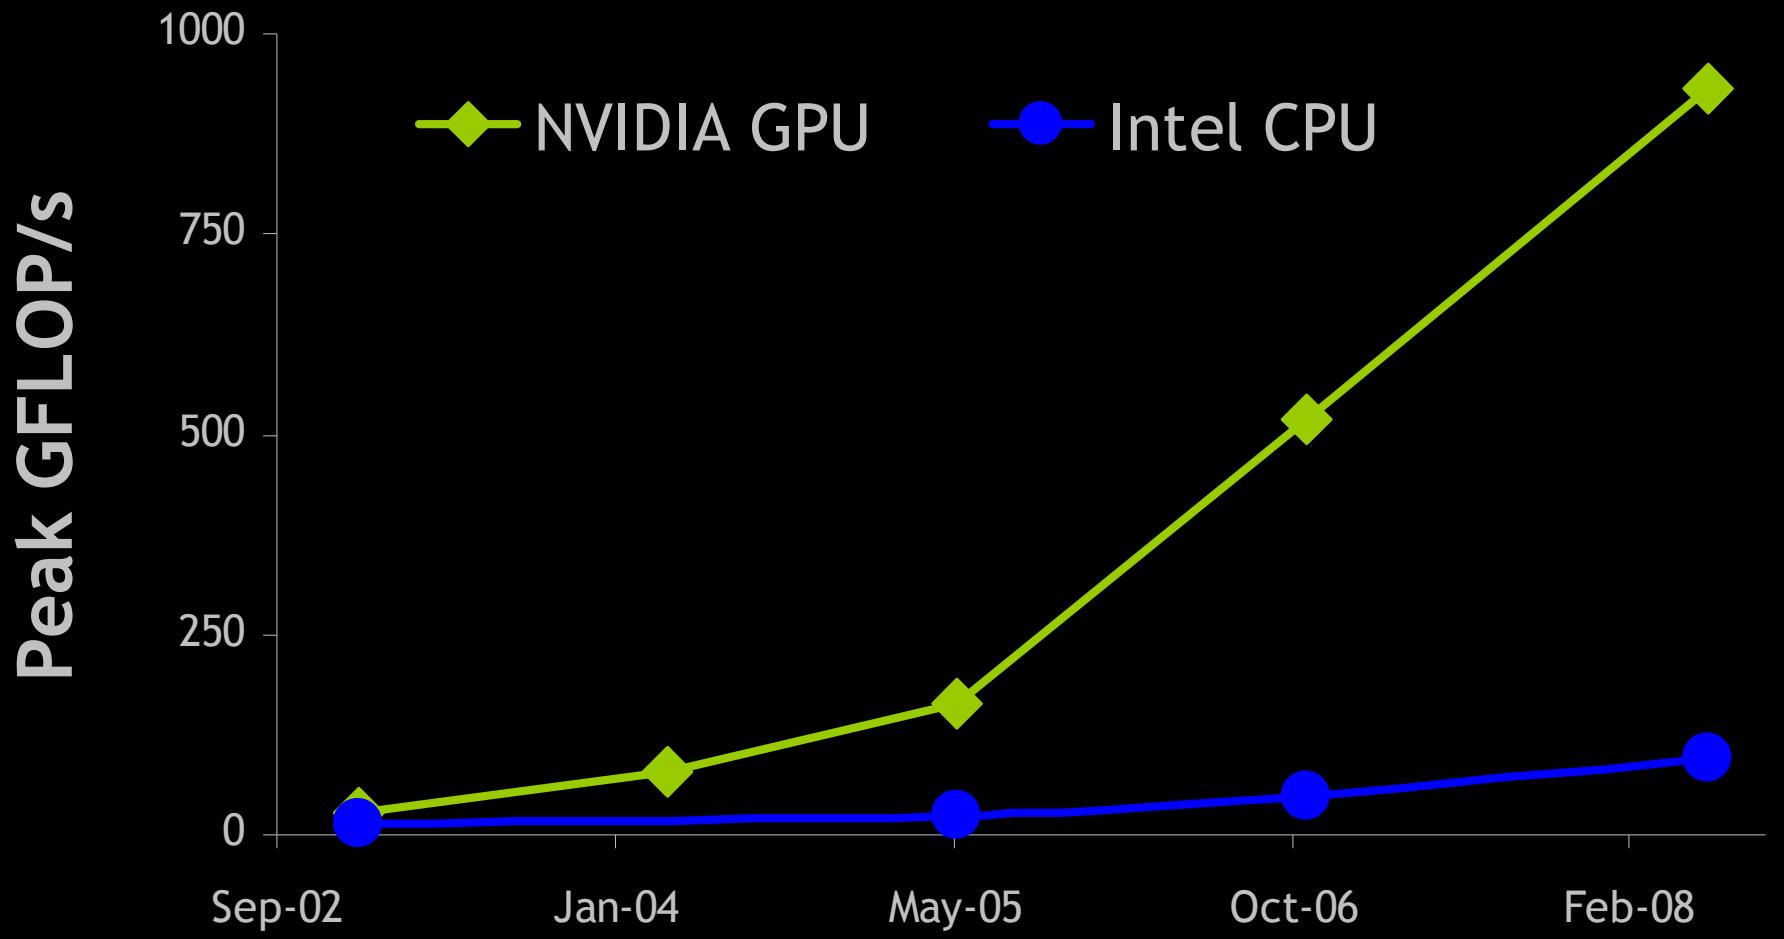
GPUs + CUDA

=

The Democratization of Parallel Computing

**Massively parallel computing has become
a commodity technology**

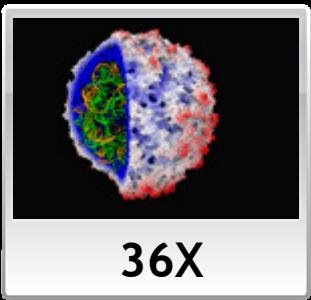
MOTIVATION



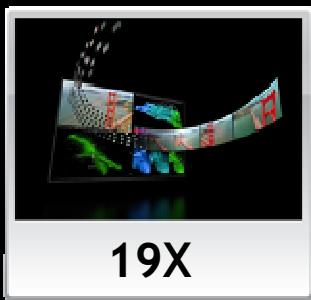
MOTIVATION



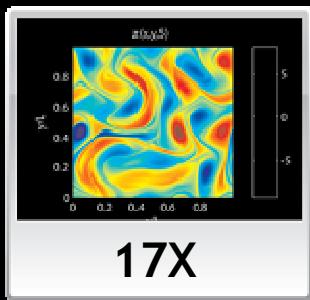
146X



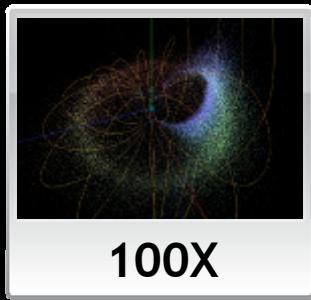
36X



19X



17X



100X

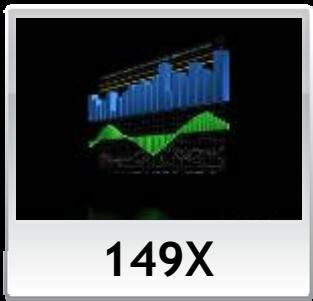
Interactive visualization of volumetric white matter connectivity

Ionic placement for molecular dynamics simulation on GPU

Transcoding HD video stream to H.264

Fluid mechanics in Matlab using .mex file CUDA function

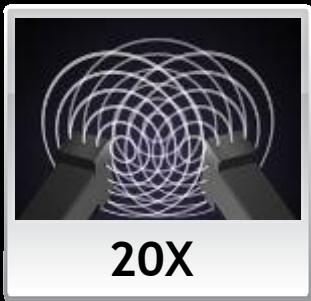
Astrophysics N-body simulation



149X



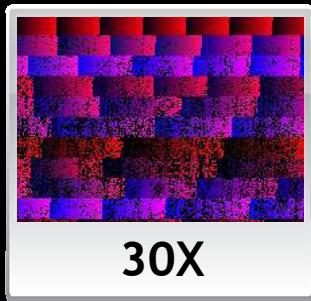
47X



20X



24X



30X

Financial simulation of LIBOR model with swaptions

GLAME@lab: an M-script API for GPU linear algebra

Ultrasound medical imaging for cancer diagnostics

Highly optimized object oriented molecular dynamics

Cmatch exact string matching to find similar proteins and gene sequences



Motivation: NVIDIA



Supercomputing Performance

- 960 cores. 4 TeraFLOPS
- 250x the performance of a desktop

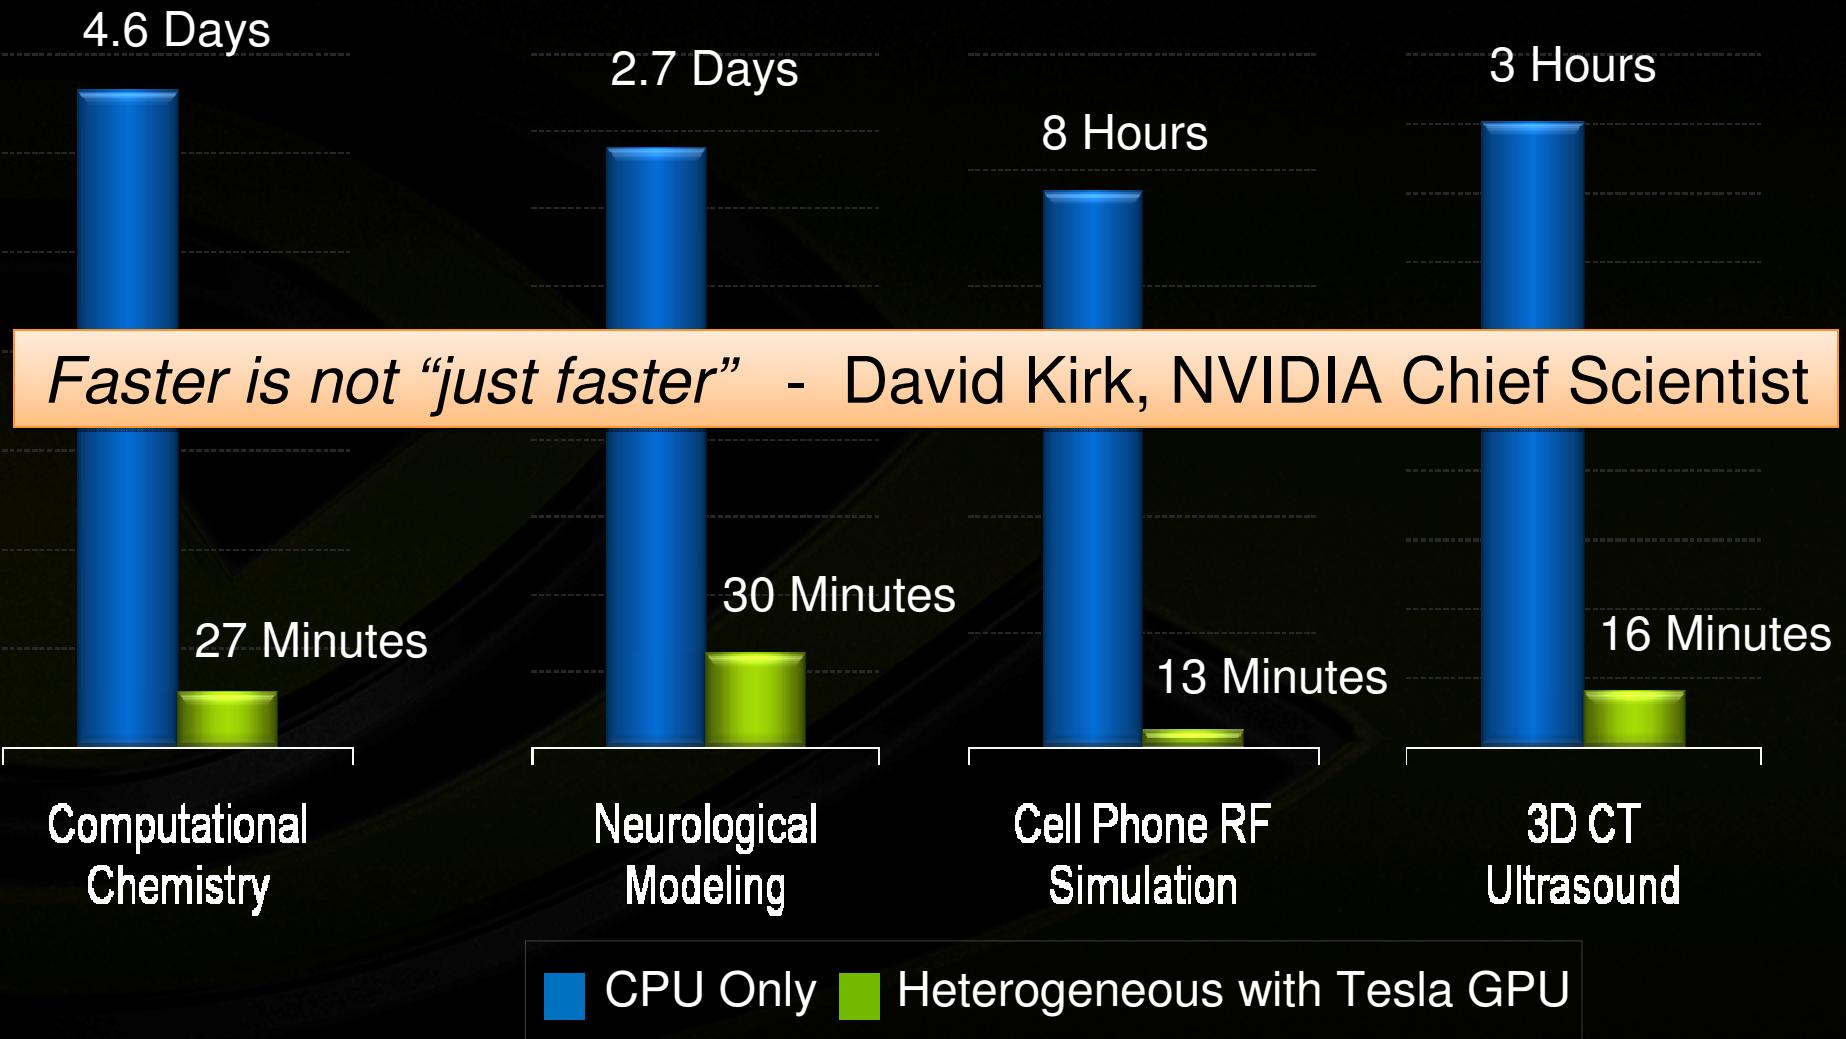
Personal

- One researcher, one supercomputer
- Plugs into standard power strip

Accessible

- Program in C for Windows, Linux
- Available now under \$10,000

Accelerating Time to Insight



CUDA: 'C' FOR PARALLELISM



```
void saxpy_serial(int n, float a, float *x, float *y)
{
    for (int i = 0; i < n; ++i)
        y[i] = a*x[i] + y[i];
}
// Invoke serial SAXPY kernel
saxpy_serial(n, 2.0, x, y);
```

Standard C Code

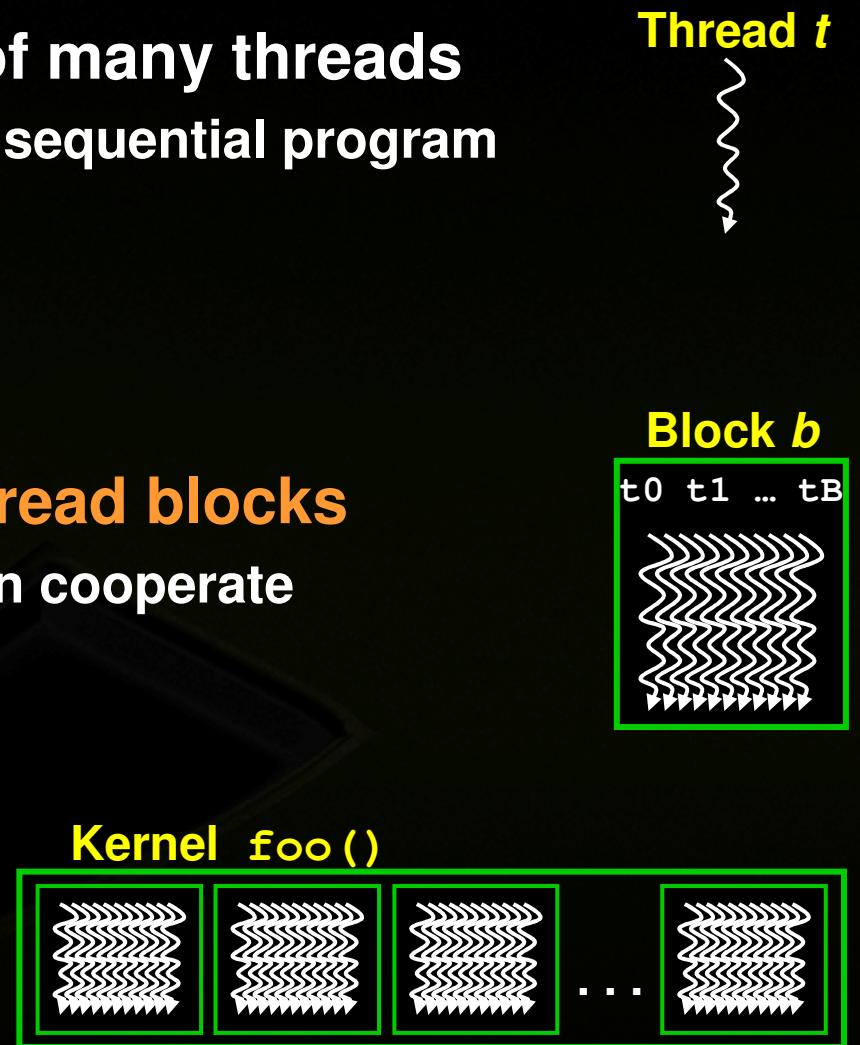
```
__global__ void saxpy_parallel(int n, float a, float *x, float *y)
{
    int i = blockIdx.x*blockDim.x + threadIdx.x;
    if (i < n)  y[i] = a*x[i] + y[i];
}
// Invoke parallel SAXPY kernel with 256 threads/block
int nbBlocks = (n + 255) / 256;
saxpy_parallel<<<nbBlocks, 256>>>(n, 2.0, x, y);
```

Parallel C Code



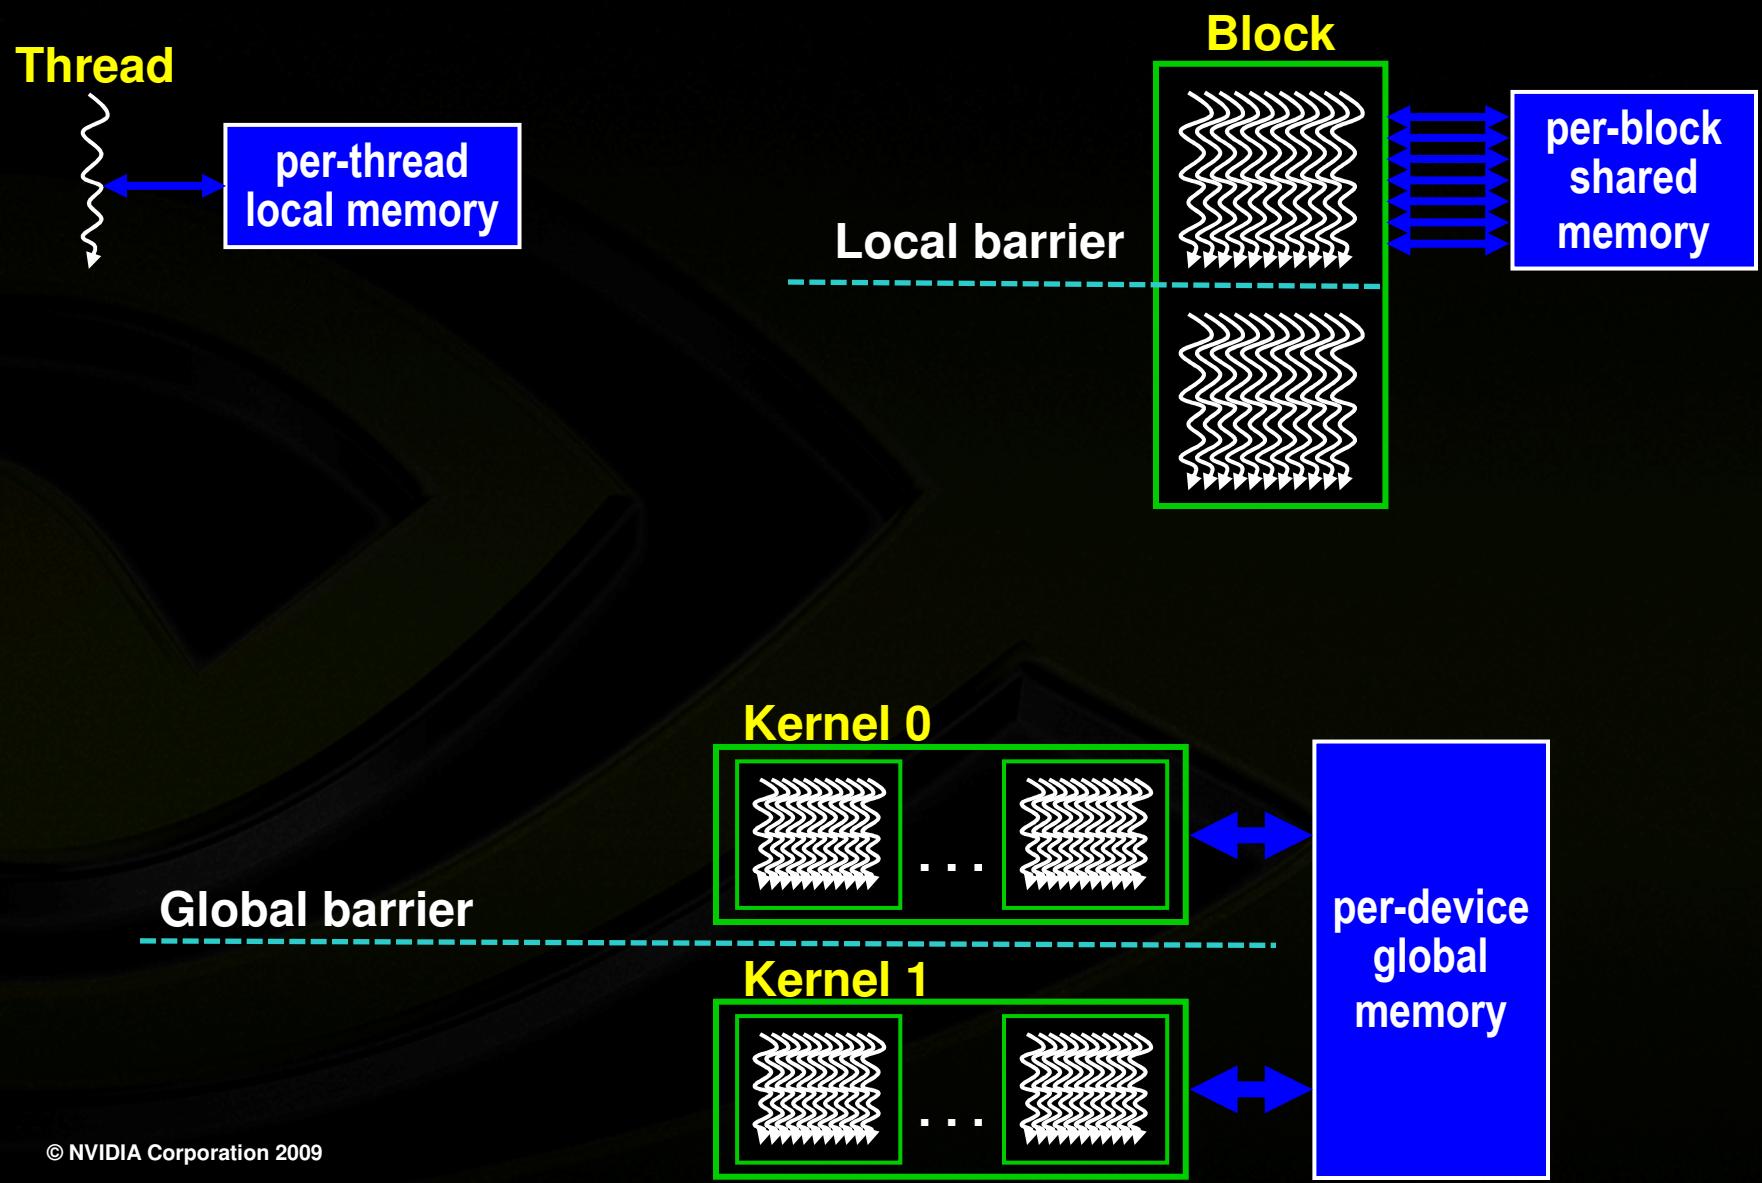
Hierarchy of concurrent threads

- Parallel **kernels** composed of many threads
 - all threads execute the same sequential program
- Threads are grouped into **thread blocks**
 - threads in the same block can cooperate
- Threads/blocks have unique IDs





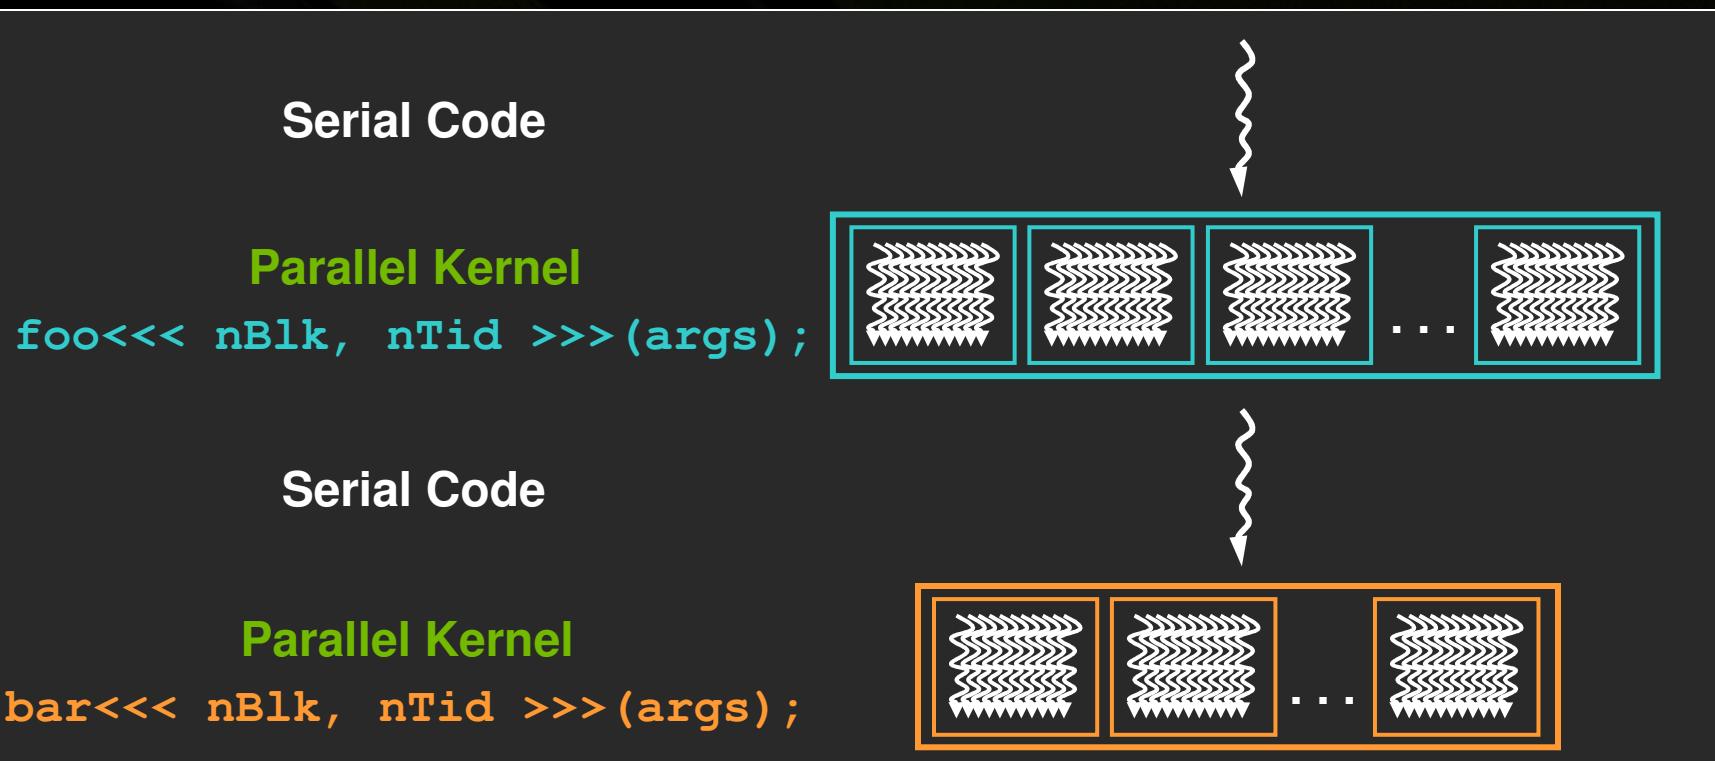
Hierarchical organization





Heterogeneous Programming

- CUDA = serial program with parallel kernels, all in C
 - Serial C code executes in a CPU thread
 - Parallel kernel C code executes in thread blocks across multiple processing elements



Thread = virtualized scalar processor



- **Independent thread of execution**
 - has its own PC, variables (registers), processor state, etc.
 - no implication about how threads are scheduled

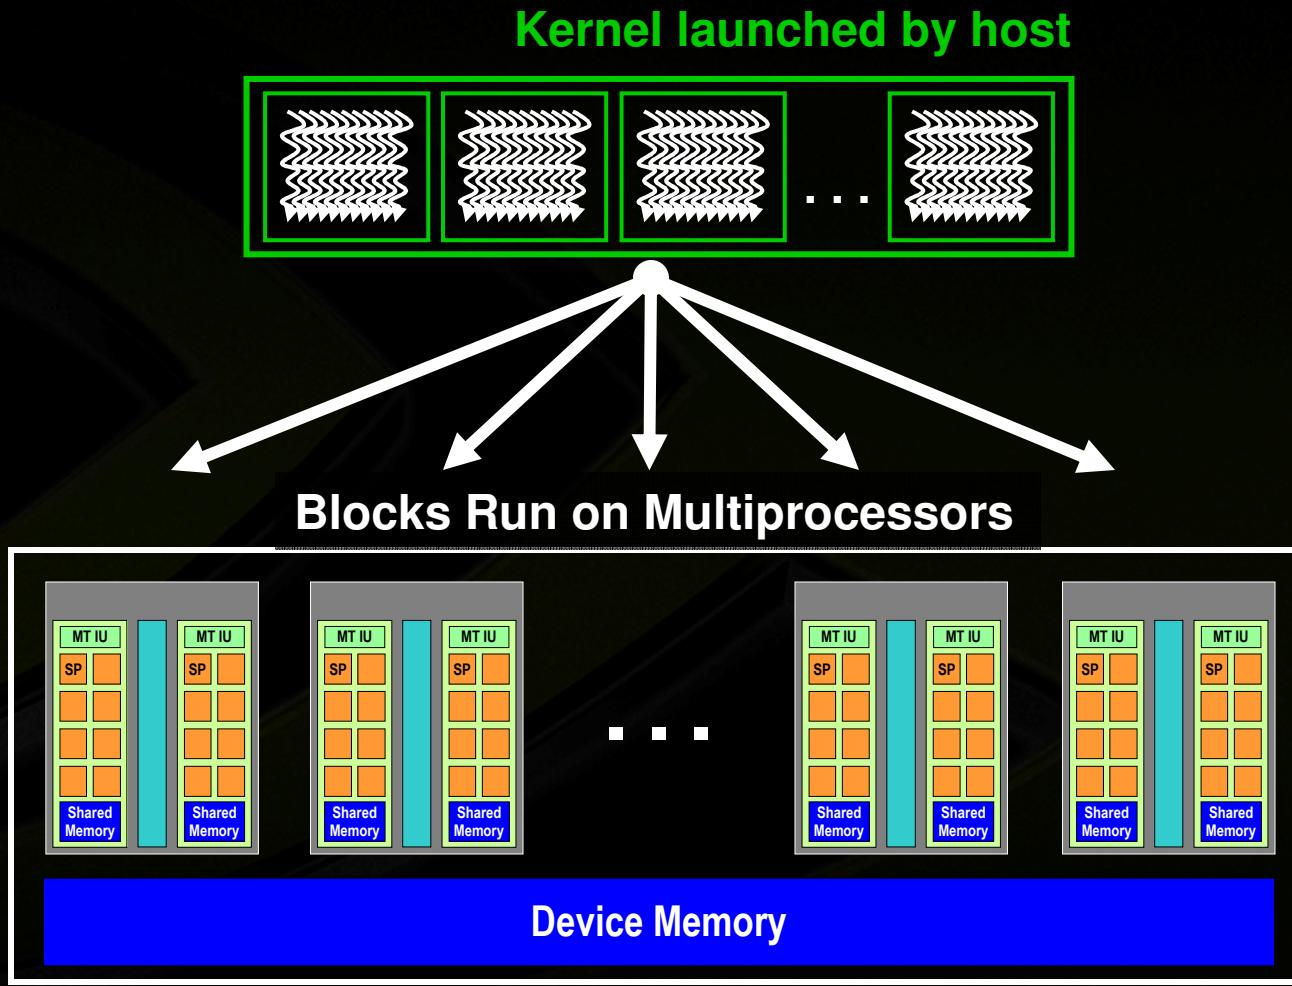
Block = virtualized multiprocessor



- Provides programmer flexibility
 - freely choose processors to fit data
 - freely customize for each kernel launch
- Thread block = a (data) parallel task
 - all blocks in kernel have the same entry point
 - but may execute any code they want
- Thread blocks of kernel must be independent tasks
 - program valid for *any interleaving* of block executions



Scalable Execution Model





Synchronization & Cooperation

- Threads within block may synchronize with **barriers**

```
... Step 1 ...
__syncthreads();
... Step 2 ...
```

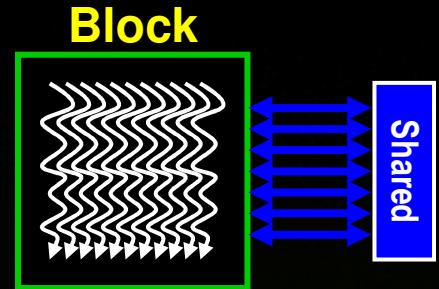
- Blocks **coordinate** via atomic memory operations
 - e.g., increment shared queue pointer with **atomicInc()**
- Implicit barrier between **dependent kernels**

```
vec_minus<<<nblocks, blksize>>>(a, b, c);
vec_dot<<<nblocks, blksize>>>(c, c);
```

Using per-block shared memory

- Variables shared across block

```
__shared__ int *begin, *end;
```



- Scratchpad memory

```
__shared__ int scratch[blocksize];
scratch[threadIdx.x] = begin[threadIdx.x];
// ... compute on scratch values ...
begin[threadIdx.x] = scratch[threadIdx.x];
```

- Communicating values between threads

```
scratch[threadIdx.x] = begin[threadIdx.x];
__syncthreads();
int left = scratch[threadIdx.x - 1];
```



Summing Up

- **CUDA = C + a few simple extensions**
 - makes it easy to start writing basic parallel programs
- **Three key abstractions:**
 1. hierarchy of parallel threads
 2. corresponding levels of synchronization
 3. corresponding memory spaces
- **Supports massive parallelism of manycore GPUs**

SOME FINAL THOUGHTS



- We should teach parallel computing in CS 1 or CS 2
- Remember: computers don't get faster, just wider
- Heapsort and mergesort
 - Both $O(n \lg n)$
 - One parallel-friendly, one not
 - Students need to understand this early

Conclusion



- **GPUs are massively parallel manycore computers**
 - Ubiquitous - most successful parallel processor in history
 - Useful - users achieve huge speedups on real problems
- **CUDA is a powerful parallel architecture and programming model**
 - Heterogeneous - mixed serial-parallel programming
 - Scalable - hierarchical thread execution model
 - Accessible – e.g. minimal but expressive changes to C
- **They provide tremendous scope for innovative research**



Questions?



Example: Vector Add w/ Host Code





Example: Vector Addition Kernel

Device Code

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__ void vecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    // Run N/256 blocks of 256 threads each
    vecAdd<<< N/256, 256>>>(d_A, d_B, d_C);
}
```



Example: Vector Addition Kernel

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__ void vecAdd(float* A, float* B, float* C)
{
    int i = threadIdx.x + blockDim.x * blockIdx.x;
    C[i] = A[i] + B[i];
}

int main()
{
    // Run N/256 blocks of 256 threads each
    vecAdd<<< N/256, 256>>>(d_A, d_B, d_C);
}
```

Host Code



Example: Host code for vecAdd

```
// allocate and initialize host (CPU) memory
float *h_A = ..., *h_B = ...;

// allocate device (GPU) memory
float *d_A, *d_B, *d_C;
cudaMalloc( (void**) &d_A, N * sizeof(float));
cudaMalloc( (void**) &d_B, N * sizeof(float));
cudaMalloc( (void**) &d_C, N * sizeof(float));

// copy host memory to device
cudaMemcpy( d_A, h_A, N * sizeof(float),
            cudaMemcpyHostToDevice) );
cudaMemcpy( d_B, h_B, N * sizeof(float),
            cudaMemcpyHostToDevice) );

// execute the kernel on N/256 blocks of 256 threads each
vecAdd<<<N/256, 256>>>(d_A, d_B, d_C);
```

**Example:
Reduction**

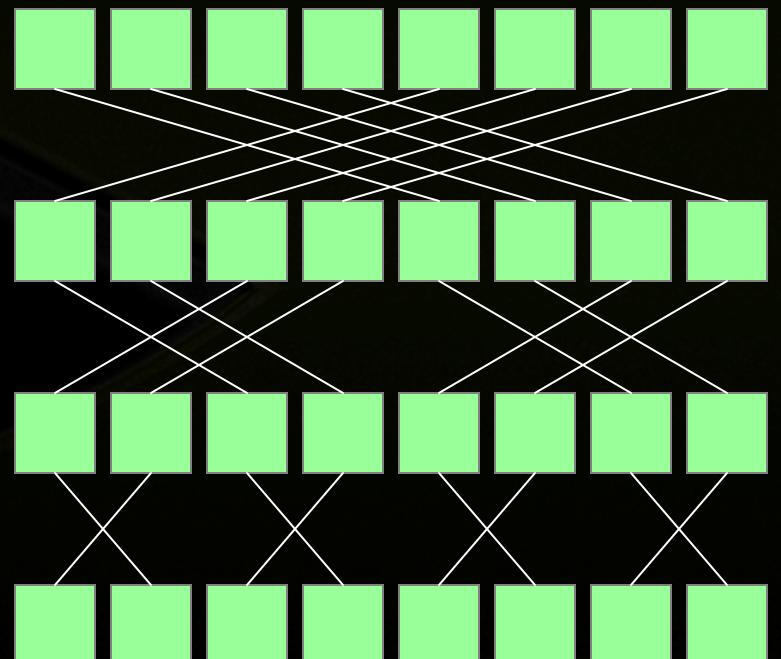
Example: Parallel Reduction

- Summing up a sequence with 1 thread:

```
int sum = 0;  
for(int i=0; i<N; ++i)    sum += x[i];
```

- Parallel reduction builds a summation tree

- each thread holds 1 element
- stepwise partial sums
- N threads need $\log N$ steps
- one possible approach:
Butterfly pattern



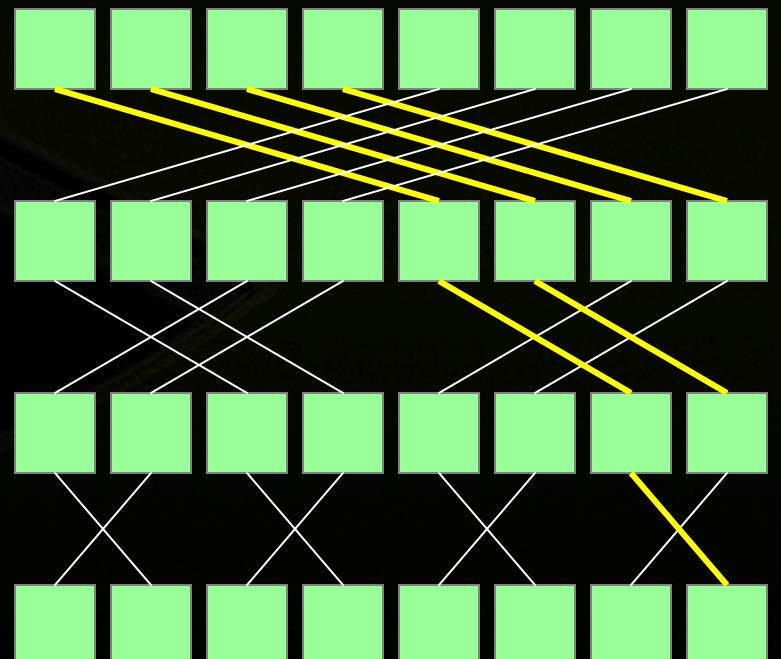
Example: Parallel Reduction

- Summing up a sequence with 1 thread:

```
int sum = 0;  
for(int i=0; i<N; ++i)    sum += x[i];
```

- Parallel reduction builds a summation tree

- each thread holds 1 element
- stepwise partial sums
- N threads need $\log N$ steps
- one possible approach:
Butterfly pattern





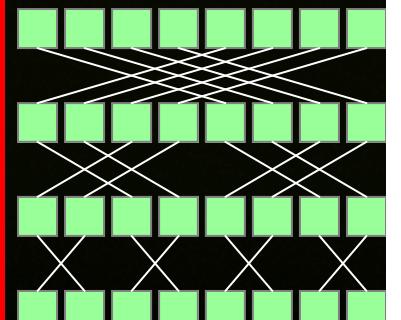
Parallel Reduction for 1 Block

```
// INPUT: Thread i holds value x_i  
int i = threadIdx.x;  
__shared__ int sum[blocksize];
```

```
// One thread per element  
sum[i] = x_i; __syncthreads();
```

```
for(int bit=blocksize/2; bit>0; bit/=2)  
{  
    int t=sum[i]+sum[i^bit]; __syncthreads();  
    sum[i]=t; __syncthreads();  
}
```

```
// OUTPUT: Every thread now holds sum in sum[i]
```





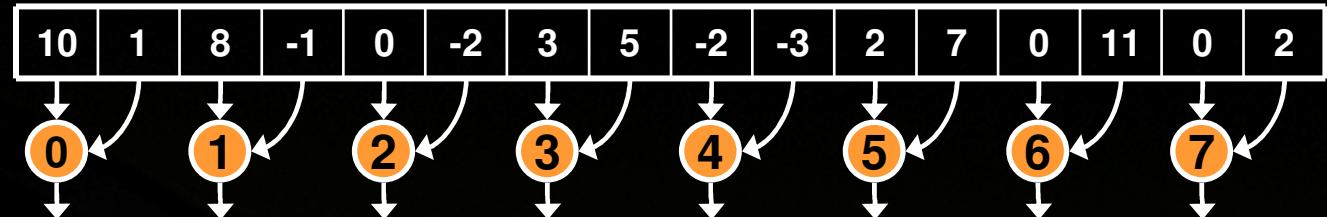
Reduction tree redux



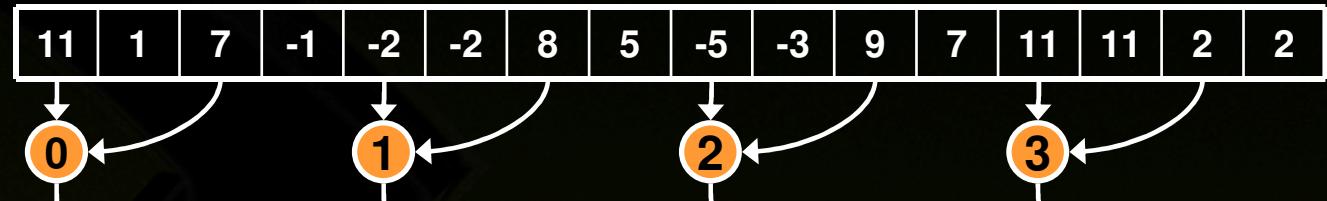


Compare to interleaved addressing:

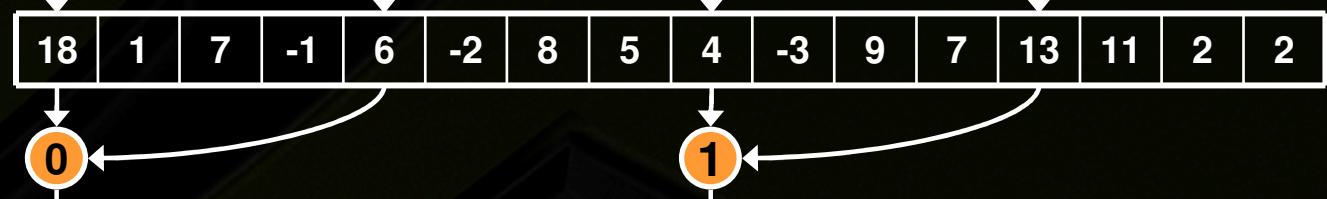
Input (shared memory)



`x[i] += x[i+8];`



`x[i] += x[i+4];`



`x[i] += x[i+2];`



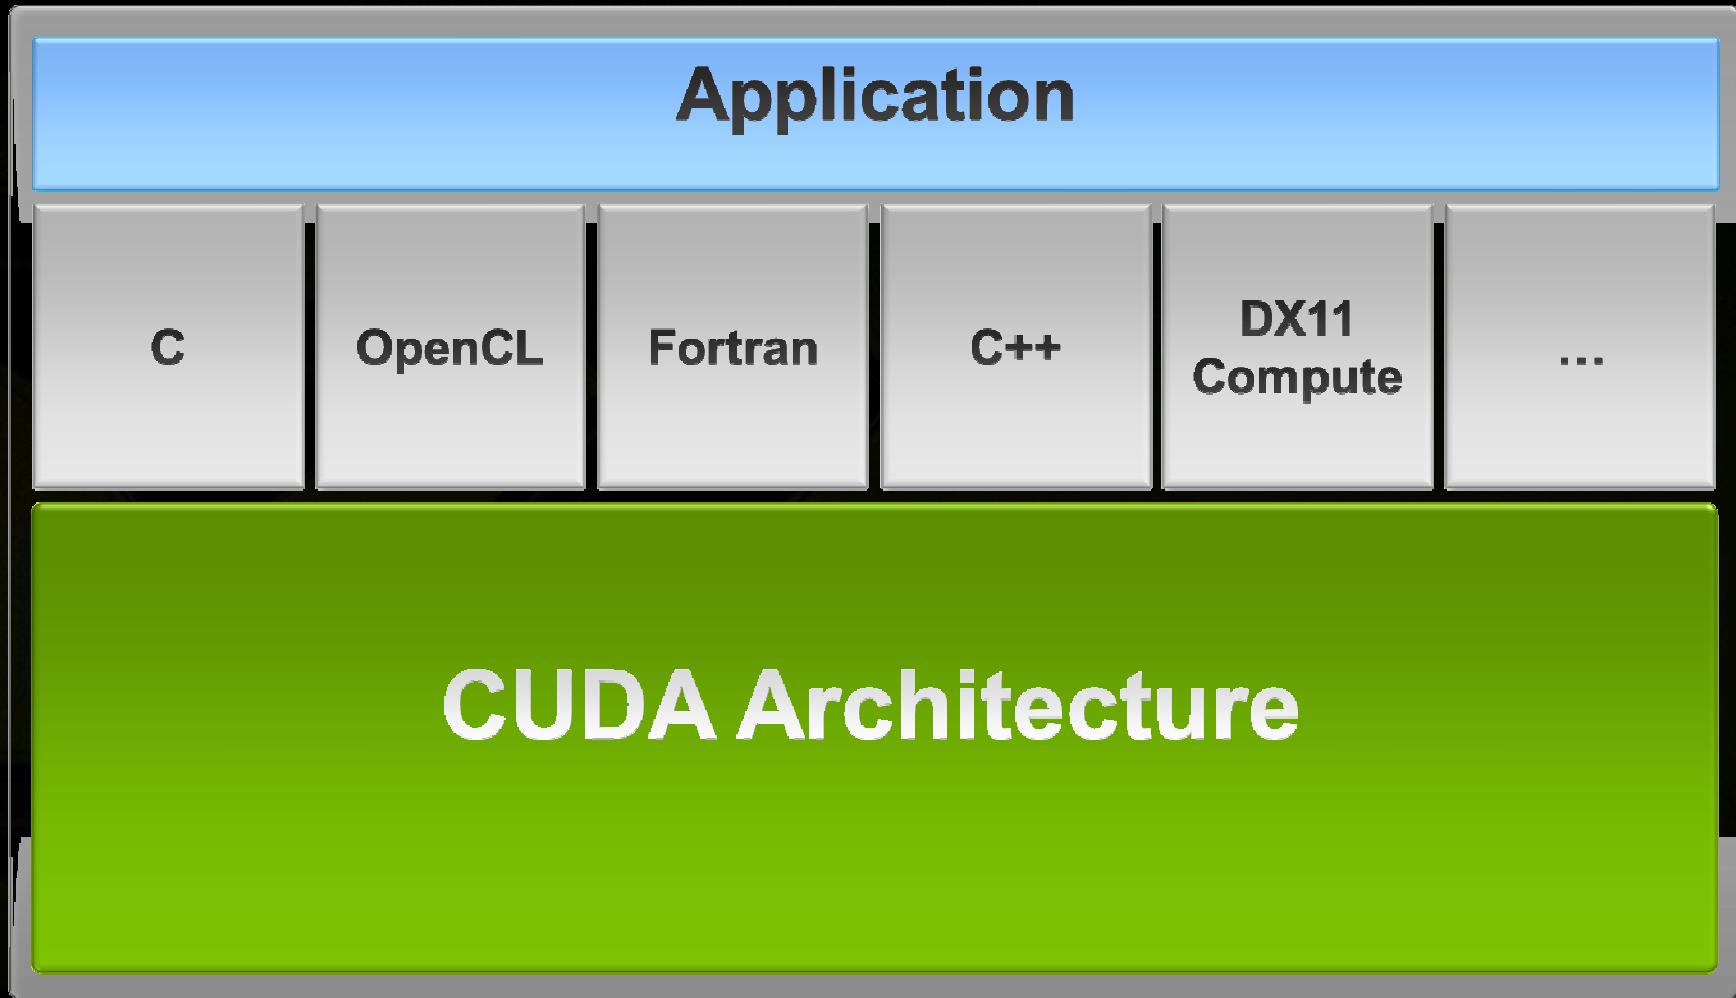
`x[i] += x[i+1];`



OpenCL

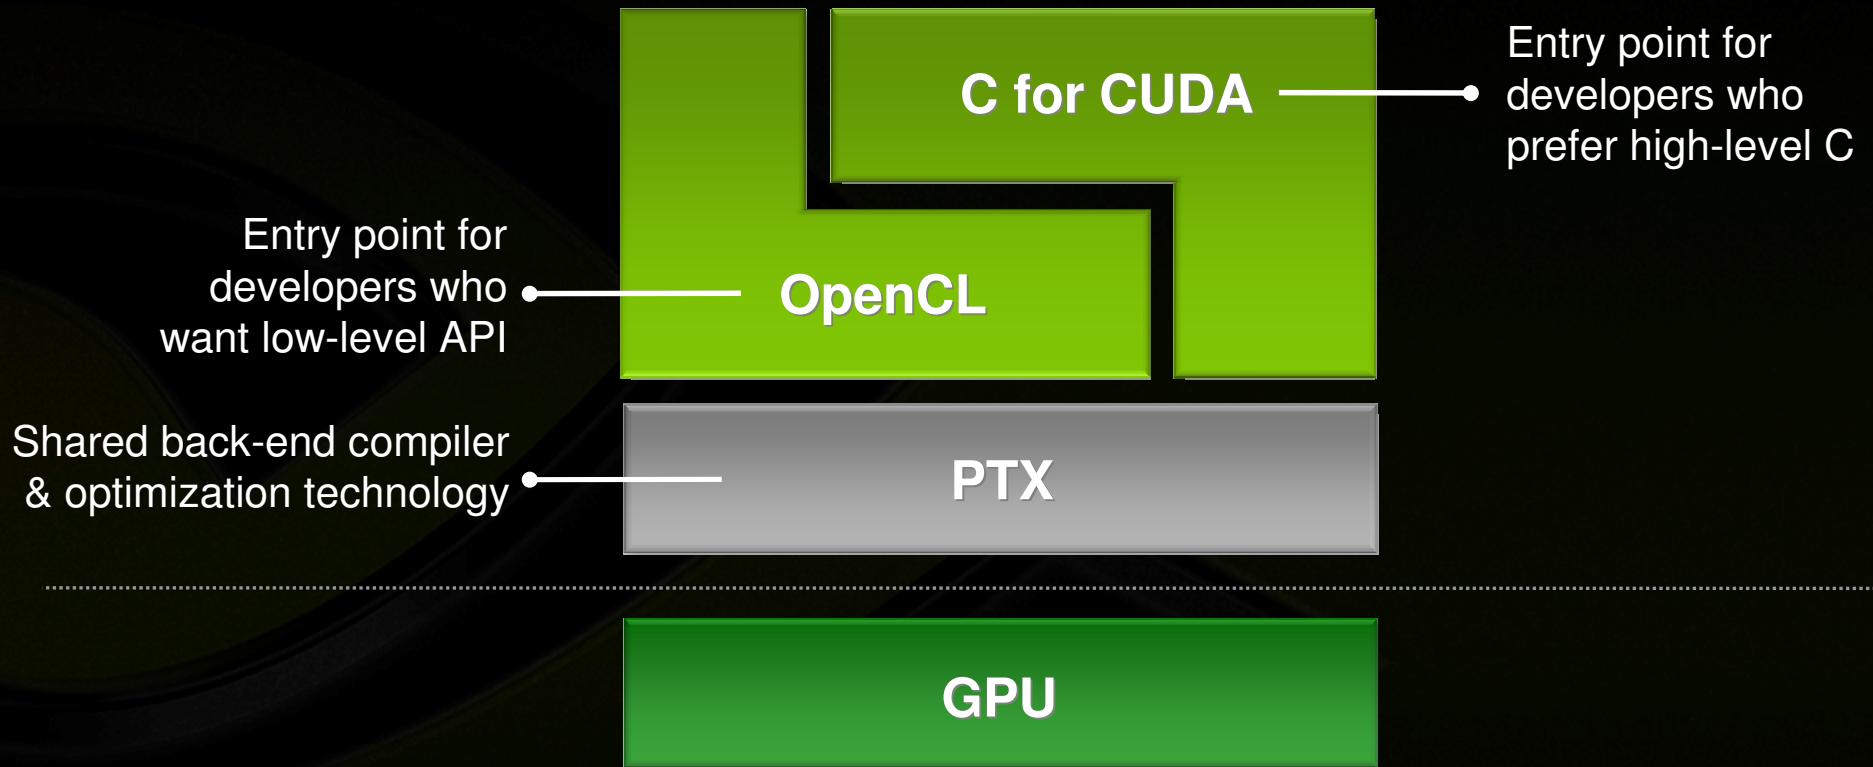


CUDA: An Architecture for Massively Parallel Computing





OpenCL vs. C for CUDA





FFT Kernel Example

OPENCL

```
__kernel void fft1D_1024 (__global float2 *in, __global float2 *out,
                           __local float *sMemx, __local float *sMemy)
{
    int tid = get_local_id(0);  int blockIdx = get_group_id(0) * 1024 + tid;
    float2 data[16];
    in = in + blockIdx;  out = out + blockIdx;
```

```
globalLoads(data, in, 64); // coalesced global reads
fftRadix16Pass(data);      // in-place radix-16 pass
twiddleFactorMul(data, tid, 1024, 0);
localShuffle(data, sMemx, sMemy, tid, (((tid & 15) * 65) + (tid >> 4)));
fftRadix16Pass(data);      // in-place radix-16 pass
twiddleFactorMul(data, tid, 64, 4); // twiddle factor multiplication
localShuffle(data, sMemx, sMemy, tid, (((tid >> 4) * 64) + (tid & 15)));
fftRadix4Pass(data);
fftRadix4Pass(data + 4); // four radix-4 function calls
fftRadix4Pass(data + 8)
fftRadix4Pass(data + 12);
globalStores(data, out, 64); // coalesced global writes
}
```

Calculate Index Load Data

FFT Kernel

C for CUDA (Written by Vasily Volkov, © UC

```
__global__ void FFT1024_device( float2 *dst, float2 *src )
{
    int tid = threadIdx.x; int iblock = blockIdx.y * gridDim.x + blockIdx.x;
    int index = iblock * 1024 + tid; src += index; dst += index;
    int hi4 = tid>>4; int lo4 = tid&15; int hi2 = tid>>4; int mi2 = (tid>>2)&3; int
        lo2 = tid&3;
    float2 a[16];
    __shared__ float smem[69*16];

    load<16>( a, src, 64 );
    FFT16( a );
    twiddle<16>( a, tid, 1024 );
    int il[] = {0,1,2,3, 16,17,18,19, 32,33,34,35, 48,49,50,51};
    transpose<16>( a, &smem[lo4*65+hi4], 4, &smem[lo4*65+hi4*4], il );
    FFT4x4( a );
    twiddle4x4( a, lo4 );
    transpose4x4( a, &smem[hi2*17 + mi2*4 + lo2], 69, &smem[mi2*69*4 +
        hi2*69 + lo2*17 ], 1, 0xE );
    FFT16( a );
    store<16>( a, dst, 64 );
}
```



Different Host Code Styles

Calling a C function in nvcc

```
extern "C" void FFT1024( float2 *work, int batch )
{
    FFT1024_device<<< grid2D(batch), 64 >>>
    ( work, work );
}
```

NVIDIA's PTX layer
manages kernel
resources and execution

OpenCL API-style programming

```
// create a compute context with GPU device
context = clCreateContextFromType(CL_DEVICE_TYPE_GPU);
// create a work-queue
queue = clCreateWorkQueue(context, NULL, NULL, 0);
// allocate the buffer memory objects
memobjs[0] = clCreateBuffer(context,
CL_MEM_READ_ONLY | CL_MEM_COPY_HOST_PTR,
sizeof(float)*2*num_entries, srcA);
memobjs[1] = clCreateBuffer(context,
CL_MEM_READ_WRITE,
sizeof(float)*2*num_entries, NULL);
// create the compute program
program = clCreateProgramFromSource(context, 1,
&fft1D_1024_kernel_src, NULL);
// build the compute program executable
clBuildProgramExecutable(program, false, NULL, NULL);
// create the compute kernel
kernel = clCreateKernel(program, "fft1D_1024");
// create N-D range object with work-item dimensions
global_work_size[0] = n;
local_work_size[0] = 64;
range = clCreateNDRangeContainer(context, 0, 1,
global_work_size,
local_work_size);
// set the args values
clSetKernelArg(kernel, 0, (void *)&memobjs[0],
sizeof(cl_mem), NULL);
clSetKernelArg(kernel, 1, (void *)&memobjs[1],
sizeof(cl_mem), NULL);
clSetKernelArg(kernel, 2, NULL,
sizeof(float)*(local_work_size[0]+1)*16, NULL);
clSetKernelArg(kernel, 3, NULL,
sizeof(float)*(local_work_size[0]+1)*16, NULL);
// execute kernel
clExecuteKernel(queue, kernel, NULL, range, NULL, 0, NULL);
```

Source:
SIGGraph sneak preview
A Munshi, Apple Computer

Sparse Linear Algebra Results

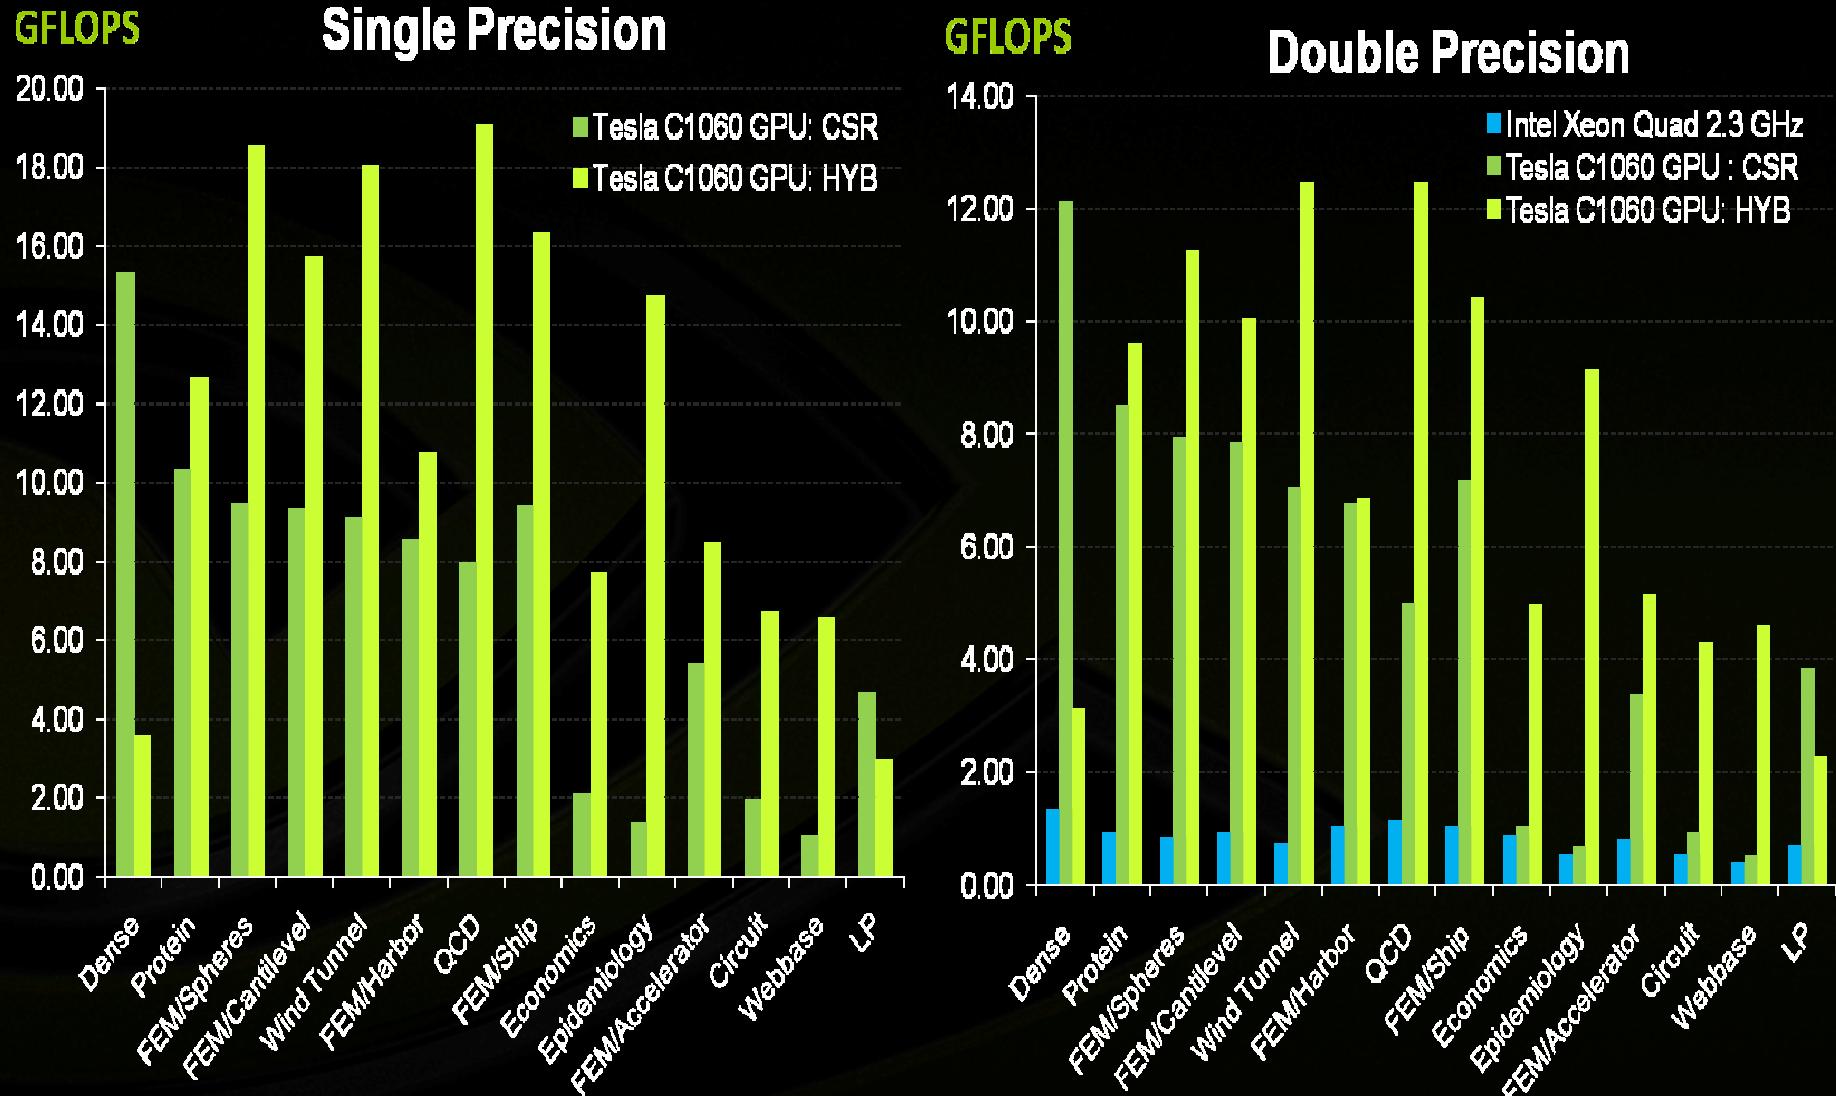


Sparse Matrix-Vector Multiplication (SpMV) on CUDA

- Experimented with several data structures
 - CSR: Compressed Sparse Row
 - HYB: Hybrid of ELLPACK (ELL) and Coordinate (COO) formats
- HYB gave best results
 - Speed of ELL with flexibility of COO
- Benchmarked against matrices from
 - *“Optimization of Sparse Matrix-Vector Multiplication on Emerging Multicore Platforms”*, S. Williams et al, Supercomputing 2007



Results: Sparse Matrix-Vector Multiplication (SpMV) on CUDA



Double Precision

T10 Double Precision Floating Point



| | |
|--|---|
| Precision | IEEE 754 |
| Rounding modes for FADD and FMUL | All 4 IEEE, round to nearest, zero, inf, -inf |
| Denormal handling | Full speed |
| NaN support | Yes |
| Overflow and Infinity support | Yes |
| Flags | No |
| FMA | Yes |
| Square root | Software with low-latency FMA-based convergence |
| Division | Software with low-latency FMA-based convergence |
| Reciprocal estimate accuracy | 24 bit |
| Reciprocal sqrt estimate accuracy | 23 bit |
| $\log_2(x)$ and 2^x estimates accuracy | 23 bit |

Double Precision Floating Point



| | NVIDIA Tesla T10 | SSE2 | Cell SPE |
|------------------------------------|---|---|---|
| Precision | IEEE 754 | IEEE 754 | IEEE 754 |
| Rounding modes for FADD and FMUL | All 4 IEEE, round to nearest, zero, inf, -inf | All 4 IEEE, round to nearest, zero, inf, -inf | All 4 IEEE, round to nearest, zero, inf, -inf |
| Denormal handling | Full speed | Supported, costs 1000's of cycles | Supported only for results, not input operands (input denormals flushed-to-0) |
| NaN support | Yes | Yes | Yes |
| Overflow and Infinity support | Yes | Yes | Yes |
| Flags | No | Yes | Yes |
| FMA | Yes | No | Yes |
| Square root | Software with low-latency FMA-based convergence | Hardware | Software only |
| Division | Software with low-latency FMA-based convergence | Hardware | Software only |
| Reciprocal estimate accuracy | 24 bit | 12 bit | 12 bit + step |
| Reciprocal sqrt estimate accuracy | 23 bit | 12 bit | 12 bit + step |
| log2(x) and 2^x estimates accuracy | 23 bit | No | No |



Single Precision Floating Point

| | G80 | SSE | IBM Altivec | Cell SPE |
|------------------------------------|------------------------------------|---|-----------------------------|-----------------------------|
| Precision | IEEE 754 | IEEE 754 | IEEE 754 | IEEE 754 |
| Rounding modes for FADD and FMUL | Round to nearest and round to zero | All 4 IEEE, round to nearest, zero, inf, -inf | Round to nearest only | Round to zero/truncate only |
| Denormal handling | Flush to zero | Supported, 1000's of cycles | Supported, 1000's of cycles | Flush to zero |
| NaN support | Yes | Yes | Yes | No |
| Overflow and Infinity support | Yes, only clamps to max norm | Yes | Yes | No, infinity |
| Flags | No | Yes | Yes | Some |
| Square root | Software only | Hardware | Software only | Software only |
| Division | Software only | Hardware | Software only | Software only |
| Reciprocal estimate accuracy | 24 bit | 12 bit | 12 bit | 12 bit |
| Reciprocal sqrt estimate accuracy | 23 bit | 12 bit | 12 bit | 12 bit |
| log2(x) and 2^x estimates accuracy | 23 bit | No | 12 bit | No |

Products

Tesla S1070 1U System



4 Teraflops¹

800 watts²

¹ single precision

² typical power

Tesla C1060 Board



957 Gigaflops¹

160 Watts²

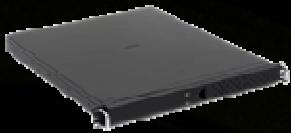
¹ single precision

² typical power

Building a 100TF datacenter



CPU 1U Server



4 CPU cores

0.07 Teraflop

\$ 2000

400 W

1429 CPU servers

\$ 3.1 M

571 KW

4 GPUs: 960 cores

4 Teraflops

\$ 8000

800 W

**25 CPU servers
25 Tesla systems**

\$ 0.31 M

27 KW

Tesla 1U System



10x lower cost

21x lower power



Tesla Personal Supercomputer



Supercomputing Performance

- Massively parallel CUDA Architecture
- 960 cores. 4 TeraFlops
- 250x the performance of a desktop

Personal

- One researcher, one supercomputer
- Plugs into standard power strip

Accessible

- Program in C for Windows, Linux
- Available now worldwide under \$10,000

C-for-CUDA SDK



Libraries:FFT, BLAS,CuDPP...
Example Source Code

Integrated CPU
and GPU C Source Code

NVIDIA C Compiler

NVIDIA Assembly
for Computing

CPU Host Code

CUDA
Driver

Debugger
Profiler

Standard C Compiler

GPU

CPU

Quotes



“ GPUs have evolved to the point where many real world applications are easily implemented on them and run significantly faster than on multi-core systems.

Future computing architectures will be hybrid systems with parallel-core GPUs working in tandem with multi-core CPUs.

”

*Jack Dongarra
Professor, University of Tennessee
Author of Linpack*



“ We've all heard ‘desktop supercomputer’ claims in the past, but this time it's for real: NVIDIA and its partners will be delivering outstanding performance and broad applicability to the mainstream marketplace.

Heterogeneous computing is what makes such a breakthrough possible.

*Burton Smith
Technical Fellow, Microsoft
Formerly, Chief Scientist at Cray*